Tollgate Client API

REST API specification · version 1.2.0 · OpenAPI 3.1.0
Tollgate: prepaid AI gateway and credit manager · Final Year Project, Information Technology University, Lahore

Rendered from `docs/api/openapi.json`, the machine-readable source of truth. Automated contract tests check that file against the running server.

Endpoints

- GET /health — Liveness check
- GET /ready — Readiness check
- GET /v1/models — List models this key may use
- POST /v1/chat/completions — Create a chat completion
- GET /v1/balance — Read your credit balance
- GET /v1/requests/{request_id} — Read one request's status and cost

Servers

- `https://api.example.com` — Replace with the base URL from your client pack.
- `http://127.0.0.1:8000` — Local development

Authentication

HTTP Bearer: `Authorization: Bearer tg_...`, a per-client virtual key issued by the operator. It is shown once and stored only as a hash.

Tollgate is a prepaid AI gateway. Your application calls Tollgate with a customer API key (`tg_...`) instead of holding AI provider keys. Every chat request reserves its worst-case cost from your credits, calls the AI provider once, then charges the exact reported usage once and returns the unused hold.

The chat endpoint is compatible with the OpenAI Chat Completions format for **text only**. Official OpenAI SDKs work by setting `base_url` to `<your base URL>/v1`.

Conventions

- **Authentication:** `Authorization: Bearer tg_...` on every `/v1/` call. `/health` and `/ready` need no key.
- **Money:** decimal strings in USD with exactly six decimal places (micro-USD), never floats.
- **Errors:** every error has the same JSON envelope `{"error": {"code", "message", "request_id", "details"?, "context"?}}`. Branch on `code`, never on `message`.
- **Request IDs:** every response carries `X-Request-ID`. Quote it when contacting support.
- **Idempotency:** send `Idempotency-Key` on chat requests so a network retry cannot generate (and charge) twice.
- **Unknown fields are rejected** with `invalid_request`. Strict types: `"256"` is not an integer.
- **Out of scope:** tools/function calling, images, audio, embeddings, `n > 1`, temperature and other sampling parameters.

Limits (operator defaults)

limit	default
requests per minute per customer	60
concurrent chat requests per customer	5
request body	262,144 bytes
messages per request	50
characters per message	8,000
characters per request (all messages)	32,000
`max_completion_tokens`	1 to 4,096 (default 256), and no more than the model's `max_output_tokens`
provider deadline	180 seconds

Request states

state	meaning
`reserved`	Credits are held; the provider has not been called yet.
`dispatched`	The provider call is in progress.
`usage_recorded`	Usage is saved; the settlement worker will charge it shortly.
`settled`	Final. The charge is in the ledger and the unused hold was released.
`released`	Final. Nothing was charged and the hold was returned.
`unresolved`	The outcome is uncertain. Credits stay held until the operator resolves it. Do not retry blindly.

Error catalog

HTTP	code	meaning	what to do
400	`content_length_mismatch`	Content-Length does not match the body that arrived.	Let your HTTP library set Content-Length.
400	`invalid_content_length`	The Content-Length header is not one non-negative integer.	Let your HTTP library set Content-Length.
401	`invalid_api_key`	The Authorization header is missing, malformed, revoked, or expired.	Send `Authorization: Bearer tg_...` with an active key. Ask the operator for a new key if it expired.
402	`insufficient_credits`	Available credits cannot cover the worst-case reservation for this request.	Top up, or lower `max_completion_tokens`. `context.required_usd` and `context.available_usd` show the gap.
403	`account_frozen`	The operator suspended new requests for this customer.	Contact the operator.
403	`model_not_allowed`	This key is not allowed to use the requested model.	Use a model listed by `GET /v1/models` for this key.
403	`spend_cap_exceeded`	The request would exceed the customer's lifetime spend cap.	Contact the operator to raise the cap.
404	`request_not_found`	No request with this ID belongs to your customer account.	Check the ID. You can only read your own requests.
404	`wrong_service`	This address serves a different role (generation vs administration).	Use the API base URL from your client pack.
408	`request_body_timeout`	The body did not arrive within the read deadline.	Retry on a stable connection.

HTTP	code	meaning	what to do
409	`duplicate_request`	The same Idempotency-Key was already used with the same body.	Do not generate again. Read `GET /v1/requests/{context.original_request_id}`.
409	`idempotency_conflict`	The same Idempotency-Key was already used with a different body or feature label.	Use a new Idempotency-Key for a different request.
409	`reservation_unavailable`	The hold was released before dispatch. No provider call was made.	Retry with a new Idempotency-Key.
413	`request_too_large`	The request body exceeds the byte limit (default 262,144 bytes).	Send a smaller body.
422	`invalid_message_order`	The provider cannot accept this sequence of roles.	Start with system (optional) then user, and alternate user/assistant.
422	`invalid_request`	A field is missing, has the wrong type, or breaks a limit. `details` lists each field.	Fix the listed fields. Do not retry unchanged.
422	`model_limit_exceeded`	Input or output size is above the model's limit (OpenAI models also need `max_completion_tokens` of at least 16).	Shorten the prompt or change `max_completion_tokens`.
422	`unsupported_model`	The model ID is not enabled on this gateway.	Use a model listed by `GET /v1/models`.
429	`concurrency_limit`	Too many chat requests are running at once for this customer (default 5).	Wait for one to finish, then retry.
429	`rate_limit_exceeded`	Too many requests from this customer in the current minute (default 60).	Wait for the `Retry-After` seconds, then retry.
500	`internal_error`	Unexpected server error. Details are hidden.	Read `GET /v1/requests/{id}` before retrying; quote `X-Request-ID` to support.
502	`provider_rejected`	The AI provider rejected or failed the call. Provider failures use code `provider_*`, `invalid_provider_*`, `incomplete_provider_stream`, `missing_usage`, `openai_*`, `anthropic_*` or `gemini_*`.	Read `GET /v1/requests/{id}`. If `released`, retry with a new Idempotency-Key. If `unresolved`, contact support.
503	`authentication_unavailable`	Key verification storage is temporarily unavailable.	Retry after a short wait.
503	`control_unavailable`	The control service is down and this node is configured to fail closed.	Retry later.
503	`price_review_required`	The operator must review this model's price before it can be used.	Retry later or contact the operator.
503	`provider_unavailable`	The AI provider could not be reached.	Read `GET /v1/requests/{id}` before retrying.
503	`storage_unavailable`	The database is temporarily unavailable.	Retry with the same Idempotency-Key.

HTTP	code	meaning	what to do
503	`warden_draining`	This API node is being taken out of service.	Retry; the load balancer will pick another node.
503	`warden_unavailable`	This API node is not registered as healthy.	Retry; the load balancer will pick another node.
503	`worker_unavailable`	Settlement is paused, so new generation is paused.	Retry later; another API node may be healthy.
504	`provider_timeout`	The provider did not finish before the deadline (default 180 s).	Read `GET /v1/requests/{id}`; credits stay held until the outcome is known.

Versioning

This document describes API version 1.2.0. Within `/v1`, changes are additive only: new optional request fields, new response fields, new error codes and new headers. Clients must ignore unknown response fields and treat unknown error codes by HTTP status. Removing or renaming anything requires a new path prefix (`/v2`) announced in advance.

GET /health

Liveness check

Confirms the process is running. It does not check the database or provider.

Authentication: None

RESPONSES

200

Process is alive.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.

application/json → HealthOut


```
{
  "status": "ok",
  "service": "Tollgate",
  "mode": "live",
  "version": "1.2.0",
  "role": "standalone"
}
```

GET /ready

Readiness check

Returns 200 when this node can accept generation: the database answers, the settlement worker is alive, and every model price is current. Load balancers should route only to ready nodes.

Authentication: None

RESPONSES

200

Ready.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.

application/json → ReadyOut

```
{
  "ready": true
}
```

503

Not ready. Do not route traffic here.

application/json → ReadyOut

GET /v1/models

List models this key may use

Returns only the models in this key's allow-list, with current prices in USD per million tokens.

Authentication: Bearer customer key (`tg_...`)

RESPONSES

200

Models available to this key.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.

application/json → ModelList


```
{
  "object": "list",
  "data": [
    {
      "id": "sim-basic",
      "object": "model",
      "owned_by": "simulator",
      "pricing_source": "test_prices",
      "price_version": "sim-test-v1",
      "reviewed_on": "",
      "input_usd_per_million": "2.000000",
      "output_usd_per_million": "8.000000",
      "cached_usd_per_million": "0.000000",
      "max_output_tokens": 4096,
      "max_input_tokens": 40000
    }
  ]
}
```


401

Error codes: `invalid_api_key`. See the error catalog in `info.description`.

application/json → ErrorResponse

404

Error codes: `wrong_service`. See the error catalog in `info.description`.

application/json → ErrorResponse

429

Error codes: `rate_limit_exceeded`. See the error catalog in `info.description`.

application/json → ErrorResponse

500

Error codes: `internal_error`. See the error catalog in `info.description`.

application/json → ErrorResponse

503

Error codes: `storage_unavailable`, `authentication_unavailable`. See the error catalog in `info.description`.

application/json → ErrorResponse

POST /v1/chat/completions

Create a chat completion

Reserves the worst-case cost, calls the provider once, and records the exact usage. The charge is settled into the ledger by the settlement worker a moment later; `GET /v1/requests/{request_id}` shows the final state.

Retry rule: retry only with the same `Idempotency-Key`. If you get `duplicate_request`, the first attempt already ran: read its status instead of generating again.

Authentication: Bearer customer key (`tg_...`)

PARAMETERS

name	in	required	rules	notes
`Idempotency-Key`	header	no	min length: `1` max length: `128` pattern: `^[A-Za-z0-9_.:-]+$`	Your unique ID for this logical request (for example a UUID). Reusing it with the same body returns `duplicate_request`; with a different body, `idempotency_conflict`.
`X-Tollgate-Feature`	header	no	min length: `1` max length: `64` pattern: `^[A-Za-z0-9_.:-]+$`	Optional label for cost reporting (for example `support-bot`). Defaults to `default`. It is part of the idempotency fingerprint.

REQUEST BODY

JSON, schema ChatCompletionRequest.

Example: Single JSON answer

```
{
  "model": "sim-basic",
  "messages": [
    {
      "role": "system",
      "content": "Answer in one paragraph."
    },
    {
      "role": "user",
      "content": "Explain an API gateway."
    }
  ],
  "max_completion_tokens": 256
}
```

Example: Streaming answer

```
{
  "model": "sim-basic",
  "messages": [
    {
      "role": "user",
      "content": "Explain an API gateway."
    }
  ],
  "max_completion_tokens": 256,
```

```
  "stream": true
}
```

RESPONSES

200

Answer generated. JSON when `stream` is false; Server-Sent Events when `stream` is true.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.
`X-Tollgate-Request-ID`	Same value as X-Request-ID. Use it with `GET /v1/requests/{request_id}`.
`Idempotency-Key`	The idempotency key in effect (yours, or the request ID if you sent none).
`X-Tollgate-Reservation-USD`	Worst-case amount held before the provider call.
`X-Tollgate-Usage-Source`	`simulated` in simulator mode, `provider_reported` in live mode.
`X-Tollgate-Warden-ID`	API node that served the request (multi-node deployments only).
`X-Tollgate-Control-Status`	`online`, or `degraded` if the control service was unreachable (multi-node only).
`X-Tollgate-Cost-USD`	Final cost of this request (JSON responses only).
`X-Tollgate-Settlement`	Request state when the response was sent, normally `usage_recorded`.

application/json → ChatCompletion

```
{
  "id": "chatcmpl-req_4f0c2b7e9a1d4c3b8e6f5a2d1c0b9a87",
  "object": "chat.completion",
  "created": 1762600000,
  "model": "sim-basic",
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "An API gateway is a single entry point..."
      },
      "finish_reason": "stop"
    }
  ],
  "usage": {
    "prompt_tokens": 14,
    "completion_tokens": 38,
    "total_tokens": 52
  }
}
```

text/event-stream: Each event is `data: <ChatCompletionChunk JSON>` followed by a blank line. The stream ends with `data: [DONE]`. If the provider fails mid-stream, a single `data: <StreamError JSON>` event is sent instead of `[DONE]` and the connection closes.

```
data: {"id":"chatcmpl-req_4f0c...","object":"chat.completion.chunk","created":1762600000,"model":"sim-
basic","choices":[{"index":0,"delta":{"content":"An API gate"},"finish_reason":null}],"usage":null}

data: {"id":"chatcmpl-req_4f0c...","object":"chat.completion.chunk","created":1762600000,"model":"sim-
basic","choices":[{"index":0,"delta":{},"finish_reason":"stop"}],"usage":null}

data: {"id":"chatcmpl-req_4f0c...","object":"chat.completion.chunk","created":1762600000,"model":"sim-
basic","choices":[],"usage":{"prompt_tokens":14,"completion_tokens":38,"total_tokens":52}}
```

```
data: [DONE]
```

400

Error codes: `invalid_content_length`, `content_length_mismatch`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

401

Error codes: `invalid_api_key`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

402

Error codes: `insufficient_credits`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

403

Error codes: `model_not_allowed`, `account_frozen`, `spend_cap_exceeded`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

404

Error codes: `wrong_service`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

408

Error codes: `request_body_timeout`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

409

Error codes: `duplicate_request`, `idempotency_conflict`, `reservation_unavailable`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

413

Error codes: `request_too_large`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

422

Error codes: `invalid_request`, `unsupported_model`, `model_limit_exceeded`, `invalid_message_order`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

429

Error codes: `rate_limit_exceeded`, `concurrency_limit`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

500

Error codes: `internal_error`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

502

Error codes: `provider_rejected`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

503

Error codes: `storage_unavailable`, `authentication_unavailable`, `price_review_required`, `provider_unavailable`, `worker_unavailable`, `warden_draining`, `warden_unavailable`, `control_unavailable`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

504

Error codes: `provider_timeout`. See the error catalog in `info.description`. Tollgate headers are present only when the failure happened after the reservation.

application/json → ErrorResponse

GET /v1/balance

Read your credit balance

`available_usd` = `balance_usd` − `reserved_usd`. `spent_usd` is lifetime settled spending.

Authentication: Bearer customer key (`tg_...`)

RESPONSES

200

Current wallet.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.

application/json → Balance


```
{
  "frozen": false,
  "balance_usd": "1.000000",
  "reserved_usd": "0.000000",
  "available_usd": "0.999668",
  "spent_usd": "0.000332",
  "spend_cap_usd": null
}
```


401

Error codes: `invalid_api_key`. See the error catalog in `info.description`.

application/json → ErrorResponse

404

Error codes: `wrong_service`. See the error catalog in `info.description`.

application/json → ErrorResponse

429

Error codes: `rate_limit_exceeded`. See the error catalog in `info.description`.

application/json → ErrorResponse

500

Error codes: `internal_error`. See the error catalog in `info.description`.

application/json → ErrorResponse

503

Error codes: `storage_unavailable`, `authentication_unavailable`. See the error catalog in `info.description`.

application/json → ErrorResponse

GET /v1/requests/{request_id}

Read one request's status and cost

Use the `X-Tollgate-Request-ID` from a chat response. You can only read your own requests.

Authentication: Bearer customer key (`tg_...`)

PARAMETERS

name	in	required	rules	notes
`request_id`	path	yes	pattern: `^req_[0-9a-f]{32}$`	

RESPONSES

200

Request metadata. Prompts and answers are never stored.

header	meaning
`X-Request-ID`	Unique ID for this HTTP request. Quote it to support.

application/json → RequestStatus


```
{
  "id": "req_4f0c2b7e9a1d4c3b8e6f5a2d1c0b9a87",
  "customer_id": "cus_3a5f0e1d2c4b6a7980f1e2d3",
  "model": "sim-basic",
  "feature": "default",
  "state": "settled",
  "price_version": "sim-test-v1",
  "input_tokens": 14,
  "output_tokens": 38,
  "error_code": null,
  "provider": "simulator",
  "provider_response_id": null,
  "warden_id": null,
  "control_degraded": false,
  "cached_tokens": 0,
  "cache_write_tokens": 0,
  "resolved_by": null,
  "created_at": 1762600000.25,
  "reserved_usd": "0.002144",
  "cost_usd": "0.000332",
  "usage_source": "simulated"
}
```


401

Error codes: `invalid_api_key`. See the error catalog in `info.description`.

application/json → ErrorResponse

404

Error codes: `wrong_service`, `request_not_found`. See the error catalog in `info.description`.

application/json → ErrorResponse

429

Error codes: `rate_limit_exceeded`. See the error catalog in `info.description`.

application/json → ErrorResponse

500

Error codes: `internal_error`. See the error catalog in `info.description`.

application/json → ErrorResponse

503

Error codes: `storage_unavailable`, `authentication_unavailable`. See the error catalog in `info.description`.

application/json → ErrorResponse

Data schemas

ChatCompletionRequest

Unknown fields are rejected.

field	type	required	rules	notes
`model`	string	yes	min length: `1` max length: `128` pattern: `^[A-Za-z0-9][A-Za-z0-9._:/-]*$`	A model ID from `GET /v1/models`.
`messages`	array of Message	yes	min items: `1` max items: `50`	
`max_completion_tokens`	integer	no	min: `1` max: `4096` default: `256`	Maximum output tokens. The legacy name `max_tokens` is accepted as an alias. The reservation is based on this value, so smaller values need less available credit.
`stream`	boolean	no	default: `false`	If true, the response is a `text/event-stream` of `ChatCompletionChunk` events ending with `data: [DONE]`.
`stream_options`	object | null	no	default: `null`	Accepted for OpenAI compatibility. The final usage chunk is always sent, with or without it.

Message

Unknown fields are rejected.

field	type	required	rules	notes
`role`	`"system"` | `"user"` | `"assistant"`	yes		
`content`	string	yes	min length: `1` max length: `8000`	

ChatCompletion

Unknown fields are rejected.

field	type	required	rules	notes
`id`	string	yes		
`object`	`"chat.completion"`	yes		
`created`	integer	yes		
`model`	string	yes		
`choices`	array of Choice	yes	min items: `1` max items: `1`	
`usage`	Usage	yes		

Choice

Unknown fields are rejected.

field	type	required	rules	notes
`index`	`0`	yes		

`message`	AssistantMessage	yes		
`finish_reason`	`"stop"` | `"length"` | `"content_filter"`	yes		

AssistantMessage

Unknown fields are rejected.

field	type	required	rules	notes
`role`	`"assistant"`	yes		
`content`	string	yes		

Usage

Unknown fields are rejected.

field	type	required	rules	notes
`prompt_tokens`	integer	yes	min: `0`	
`completion_tokens`	integer	yes	min: `0`	
`total_tokens`	integer	yes	min: `0`	
`prompt_tokens_details`	PromptTokenDetails | null	no	default: `null`	
`completion_tokens_details`	CompletionTokenDetails | null	no	default: `null`	
`cache_write_tokens`	integer | null	no	default: `null` min: `0`	

PromptTokenDetails

Unknown fields are rejected.

field	type	required	rules	notes
`cached_tokens`	integer	yes	min: `0`	

CompletionTokenDetails

Unknown fields are rejected.

field	type	required	rules	notes
`reasoning_tokens`	integer	yes	min: `0`	

ChatCompletionChunk

Unknown fields are rejected.

field	type	required	rules	notes
`id`	string	yes		
`object`	`"chat.completion.chunk"`	yes		
`created`	integer	yes		
`model`	string	yes		
`choices`	array of ChunkChoice	yes	max items: `1`	
`usage`	Usage | null	yes		

ChunkChoice

Unknown fields are rejected.

field	type	required	rules	notes
`index`	`0`	yes		
`delta`	Delta	yes		
`finish_reason`	`"stop"` | `"length"` | `"content_filter"` | null	yes		

Delta

Unknown fields are rejected.

field	type	required	rules	notes
`content`	string | null	no	default: `null`	

StreamError

Unknown fields are rejected.

field	type	required	rules	notes
`error`	StreamErrorBody	yes		

StreamErrorBody

Unknown fields are rejected.

field	type	required	rules	notes
`code`	string	yes		
`message`	string	yes		
`request_id`	string	yes		

ModelList

Unknown fields are rejected.

field	type	required	rules	notes
`object`	`"list"`	yes		
`data`	array of ModelOut	yes		

ModelOut

Unknown fields are rejected.

field	type	required	rules	notes
`id`	string	yes		
`object`	`"model"`	yes		
`owned_by`	`"simulator"` | `"openai"` | `"anthropic"` | `"gemini"`	yes		
`pricing_source`	`"test_prices"` | `"free_tier_assumed"` | `"configured_provider_rates"`	yes		
`price_version`	string	yes		
`reviewed_on`	string	yes		

`input_usd_per_million`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`output_usd_per_million`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`cached_usd_per_million`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`max_output_tokens`	integer	yes		
`max_input_tokens`	integer	yes		

Balance

Unknown fields are rejected.

field	type	required	rules	notes
`frozen`	boolean	yes		
`balance_usd`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`reserved_usd`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`available_usd`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`spent_usd`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`spend_cap_usd`	string | null	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	

RequestStatus

Unknown fields are rejected.

field	type	required	rules	notes
`id`	string	yes		
`customer_id`	string	yes		
`model`	string	yes		
`feature`	string	yes		
`state`	`"reserved"` | `"dispatched"` | `"usage_recorded"` | `"settled"` | `"released"` | `"unresolved"`	yes		
`price_version`	string	yes		
`input_tokens`	integer | null	yes		
`output_tokens`	integer | null	yes		
`error_code`	string | null	yes		
`provider`	string | null	yes		
`provider_response_id`	string | null	yes		
`warden_id`	string | null	yes		
`control_degraded`	boolean | integer	yes		
`cached_tokens`	integer | null	yes		
`cache_write_tokens`	integer | null	yes		
`resolved_by`	string | null	yes		
`created_at`	number	yes		
`reserved_usd`	string	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	

`cost_usd`	string | null	yes	pattern: `^-?[0-9]+\.[0-9]{6}$`	
`usage_source`	`"simulated"` | `"provider_reported"`	yes		

ErrorResponse

Unknown fields are rejected.

field	type	required	rules	notes
`error`	ErrorBody	yes		

ErrorBody

Unknown fields are rejected.

field	type	required	rules	notes
`code`	string	yes		
`message`	string	yes		
`request_id`	string	yes		
`details`	array of ErrorDetail | null	no	default: `null` max items: `10`	
`context`	object | null	no	default: `null`	

HealthOut

Unknown fields are rejected.

field	type	required	rules	notes
`status`	`"ok"`	yes		
`service`	string	yes		
`mode`	`"simulator"` | `"live"`	yes		
`version`	string	yes		
`role`	`"standalone"` | `"control"` | `"warden"`	yes		

ReadyOut

Unknown fields are rejected.

field	type	required	rules	notes
`ready`	boolean	yes		

ErrorDetail

Unknown fields are rejected.

field	type	required	rules	notes
`field`	string	yes		
`message`	string	yes		
`type`	string	yes		